



Kubernetes (K8s)

Background.

- An open-source system for automating the deployment, scaling and management of containerized applications.
 - A container orchestration platform.
- Originally announced by Google in 2014.
- In 2015, Google donated it to the Cloud Native Computing Foundation (CNCF).
- Kubernetes is the Greek word for helmsman – a person that steers a ship that carries containers of goods.

The problems it solves.

- Problem 1: An application is comprised of multiple containers running on a cluster of nodes (on-premise server, Cloud VMs e.g. AWS EC2). A container or an entire node crashes.
 - K8s provides monitoring and self-healing, to ensure high availability.
- Problem 2: A container experiences a spike in traffic resulting in increased latency.
 - K8s provides auto-scaling and load balancing. It replicates the container on the node and/or across the cluster of nodes and load balances the requests between them. It scales up and down based on demand.

The problems it solves.

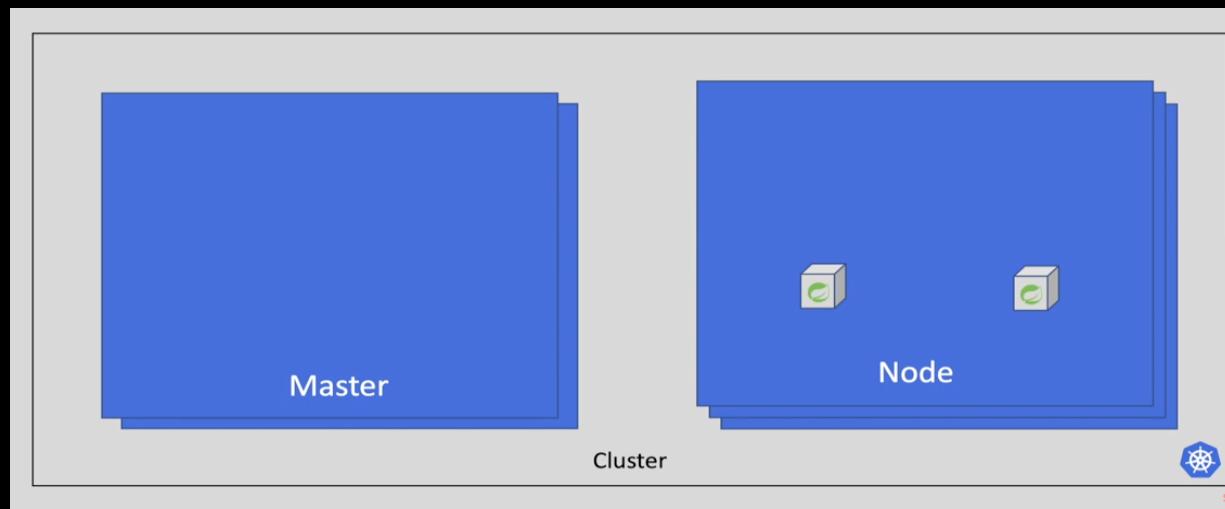
- Problem 3: You regularly upgrade an image and want to apply the change to a fleet of running containers without effecting downtime.
 - K8s provides rolling deployments - delete and restart each container, one by one rather than all at once.
 - Also supports canary deployments.
- Problem 4: I want to upgrade my entire application from v1 to v2.
 - K8s provides automatic rollout and rollback.

The problems it solves.

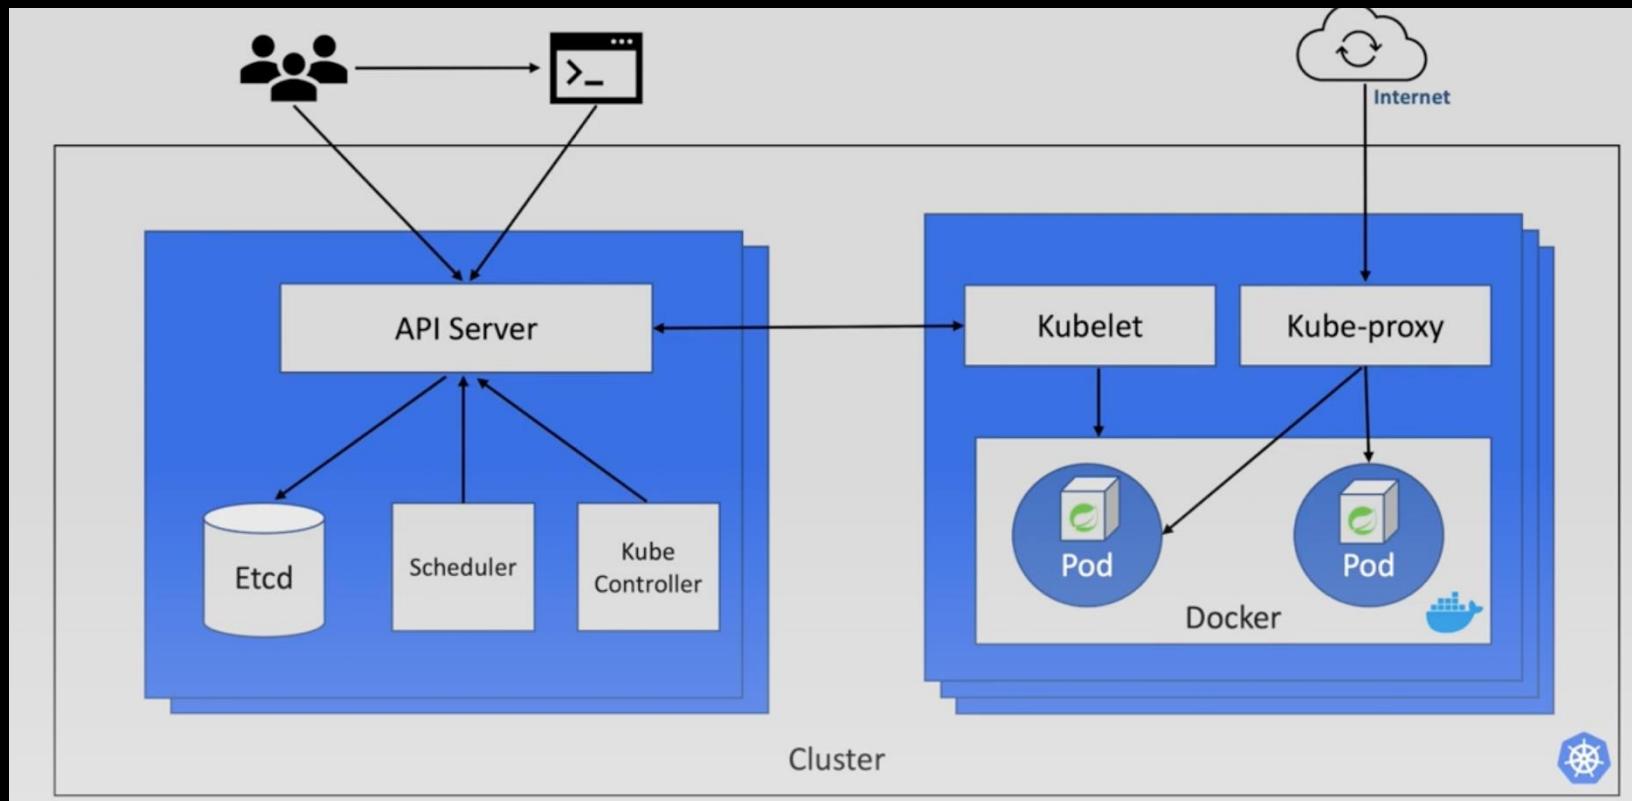
- **Problem 5: I want to reconfigure my application's without rebuilding images.**
- K8s provides secrets and configuration mappings to safeguard sensitive data and avoid unnecessary rebuilding of images.

K8s Architecture

- K8s is installed over a set of nodes (VMs)
- Worker nodes host containers - the Data plane
- Master nodes (Control plane) manages the workers – More than one master for fault tolerance and high availability.
- The set of masters and workers is called a cluster

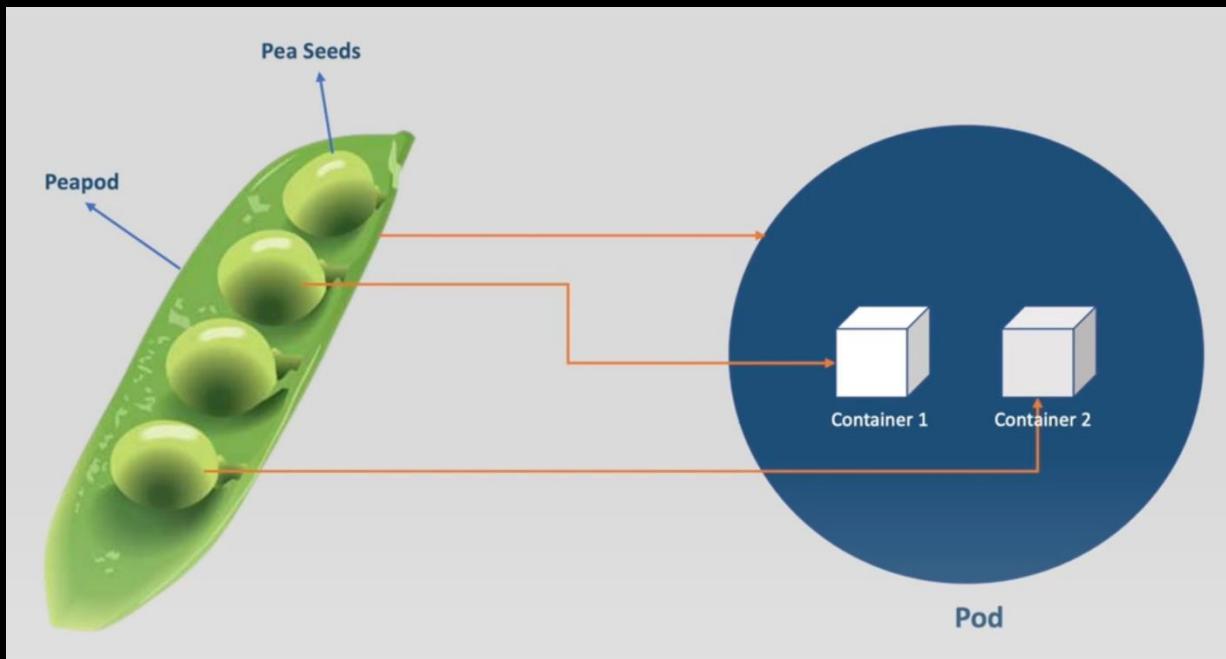


K8s Architecture




All about Pods

- The smallest unit of deployment is a pod.
- It encapsulates a group of containers, that share the same localhost network and storage.
- A pod is assigned a private IP address, but it's dynamic.



All about Pods

- When a pair of containers are tightly coupled, they must reside on the same worker node. The pod construct ensures this is satisfied.



Pod basics

- The containers in a pod share the same network (localhost), but use different ports.
- Deleting a pod will delete all its containers.
- To scale an app, we replicate its pod, where each one has its own group of containers.
- IP addressing:
 - A pod is assigned a private IP address.
 - On scale up, each pod replica is assigned its own IP.
 - Pod IPs are dynamic.

Pod basics

```
> oc run nginx-pod --image=nginx:latest  
pod/nginx-pod created
```

```
> oc get pods  
NAME      READY  STATUS    RESTARTS   AGE  
nginx-pod  1/1    Running  0          12s
```

- Creating Kubernetes resources (e.g. pods) from the command line is cumbersome; instead, we use declarative code files for better maintainability.
 - Files are termed manifests – yaml or json options.

Pod basics - Manifest files.

```
› oc apply -f 01-nginx-pod.yaml  
pod/nginx-pod1 created
```

```
› oc get pods  
NAME      READY  STATUS   RESTARTS  AGE  
nginx-pod  1/1    Running  0        3d5h  
nginx-pod1  1/1    Running  0        16s
```

```
› oc get pods -l team=integration  (Pod labels)  
NAME      READY  STATUS   RESTARTS  AGE  
nginx-pod1  1/1    Running  0        2m7s
```

Pod basics - Attributes

- The etcd stores lots of attributes about pods

```
> oc get pod nginx-pod1 -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE
READINESS GATES							
nginx-pod1	1/1	Running	0	13m	10.130.1.156	worker1	<none>

- Use describe command to get all etcd info:

```
> oc describe pod nginx-pod1
```

- Response also shows the ‘events’ that occurred for a pod.

Pod basics - Debugging

- Three main options:

1. Port-forwarding is possible with cluster access:

› oc port-forward nginx-pod1 3000:80

Forwarding from 127.0.0.1:3000 -> 80

Forwarding from [::1]:3000 -> 80

2. Check container's logs:

› oc logs nginx-pod1

3. Open a terminal shell inside the container:

› oc exec -it nginx-pod1 -- /bin/bash

root@nginx-pod1:/#

ReplicaSets

- For high availability of an app (pod), K8s can create multiple copies in the cluster, called replica sets.

```
› oc apply -f 02-nginx-replica.yaml
```

```
replicaset.apps/nginx-replica created
```

```
› oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replica-dmfvn	0/1	ContainerCreating	0	7s
nginx-replica-wc7j9	0/1	ContainerCreating	0	7s

```
› oc get rs (rs - replicaset)
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-replica	2	2	2	14s

Self-healing

```
> oc delete pod nginx-replica-dmfvn
```

```
pod "nginx-replica-dmfvn" deleted
```

```
> oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replica-kkhs2	0/1	ContainerCreating	0	3s
nginx-replica-wc7j9	1/1	Running	0	7m36s

- The new pod's IP address is different to the one it replaced.
- A replica set's `spec.selector` property determines the pods it controls.
 - A pod's labels match the replica set's selector

Deployments

- Deployments provide extra features over replica sets:
 - Rollout and rollback;
 - Deployment strategy.
 - Etc
- It automatically creates a replica set resource.
- The resource hierarchy:
 - A deployment manages a replica set.
 - A replica set manages a set of pods, based on matching the replica set's selectors with a pod's labels.
 - A pod manages its containers.

Deployments

```
› oc apply -f 03-nginx-deployment.yaml
```

```
deployment.apps/nginx-depl created
```

```
› oc get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-depl	1/2	2	1	16s

```
› oc get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-depl-7569b665c5	2	2	2	20s

```
› oc get po
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-depl-7569b665c5-dvxrj	1/1	Running	0	25s
nginx-depl-7569b665c5-nd77g	1/1	Running	0	25s

Deployment Updates

- Each Deployment update creates a new replica set

```
› oc apply -f 03-nginx-deployment.yaml. (After update)
```

```
deployment.apps/nginx-depl configured
```

```
› oc get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-depl-588f77bb97	2	2	1	15s
nginx-depl-7569b665c5	1	1	1	16m

```
› oc get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-depl-588f77bb97	2	2	2	30s
nginx-depl-7569b665c5	0	0	0	16m

Deployment Rollout

```
> oc apply -f 03-nginx-deployment.yaml (1st deploy)
```

```
deployment.apps/nginx-depl created
```

```
> oc rollout history deployment nginx-depl
```

```
deployment.apps/nginx-depl
```

```
REVISION CHANGE-CAUSE
```

```
1 Initial deployment
```

- Suppose we edit and apply the YAML twice:

```
> oc rollout history deployment nginx-depl
```

```
deployment.apps/nginx-depl
```

```
REVISION CHANGE-CAUSE
```

```
1 Initial deployment
```

```
2 Change to Nginx 1.21.5
```

```
3 Change to port 82
```

Deployment Rollback


- Rollback – Undo a Deployment rollout.

```
> oc rollout undo deployment nginx-depl  
    --to-revision=2  
deployment.apps/nginx-depl rolled back
```

```
> oc rollout history deployment nginx-depl  
deployment.apps/nginx-depl  
REVISION CHANGE-CAUSE  
1   Initial deployment  
3   Change to port 82  
4   Change to Nginx 1.21.5
```

Services

- How do pods communicate with each other?
 - Pods are assigned IP addresses, BUT these are NOT static and are not accessible outside the cluster.
 - Most rollouts and rollback cause pods to be recreated, resulting in new IP address assignments.
- Solution: Services.



Services

- Services abstract a set of pods (typically a replica set).
 - They're assigned a static IP address.
 - Unlike pods, a service is not assigned to a node.
 - Services load balance the workload between its pods.
 - Service Types:
 1. ClusterIP (Accessible inside the cluster only)
 2. NodePort
 3. LoadBalancer
 4. Ingress

ClusterIP Service creation

```
> oc apply -f 05-nginx-service.yaml
```

```
service/nginx-service created
```

```
> oc get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx-service	ClusterIP	172.30.250.186	<none>		
		8081/TCP	40s		

```
> oc get endpoints nginx-service
```

```
NAME      ENDPOINTS          AGE
```

```
nginx-service  10.130.0.75:80,10.131.1.91:80  2m10s
```

- These are the IPs (and ports) of its associated pods

ClusterIP Service access

- Service IPs are private (only accessible inside the cluster).

```
› curl 172.30.250.186:8081
```

```
^C
```

```
› oc exec -it nginx-depl-7569b665c5-2mfzg -- sh
```

```
# curl 172.30.250.186:8081
```

```
<!DOCTYPE html><html>...this works ...</html>
```

```
# curl http://nginx-service:8081
```

```
<!DOCTYPE html><html>.....</html>
```

- K8s support domain name resolution and service discovery.
- Port-forwarding with services is also allowed.

Services – Load balancing

- A service load balances the traffic to the pods

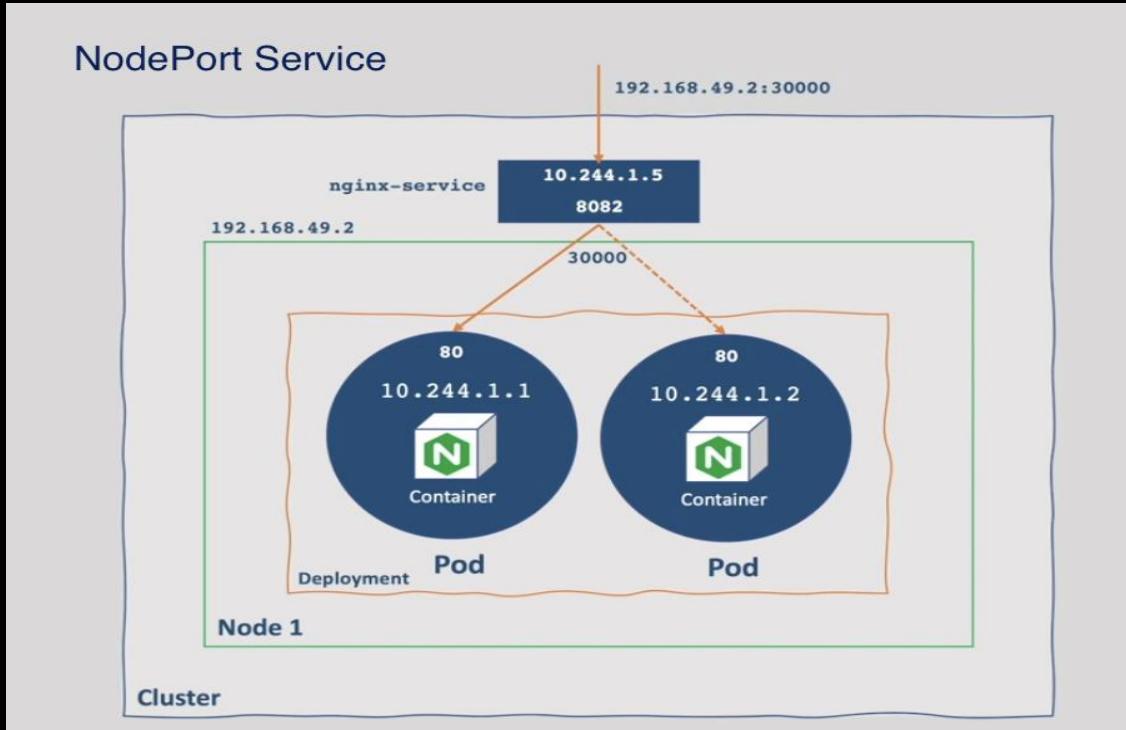
```
> oc exec -it nginx-depl-756965c5-2mfzg -- bin/bash
```

```
# i=1
while [ $i -le 20 ]
do
curl nginx-service:8081;
i=$(( $i + 1 ))
done
```

- Monitor the pods by streaming their logs:

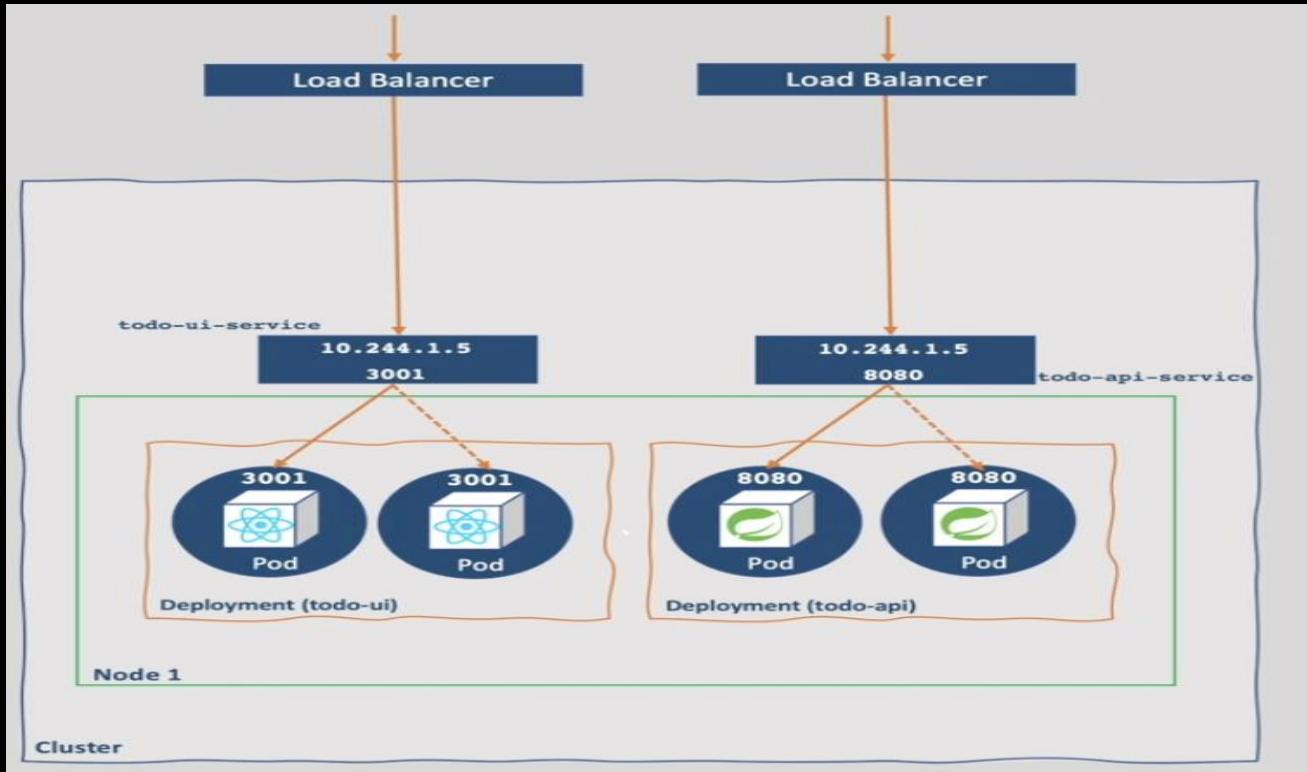
```
> oc logs nginx-depl-7569b665c5-2mfzg -f
```

NodePort Service



- Disadv: Insecure - Opens a worker's port – 1 port per service. Unstable – Node IPs change on rebooting.

LoadBalancer Service



- Disadv: Must configure an external Load Balancer, e.g. AWS Load Balancer. Costly, as each NP service needs its own LB.

Ingress Service

http://todo.com http://todo.com/api

Load Balancer or
Proxy Server

The diagram illustrates the flow of traffic through an Ingress Service. External requests for `http://todo.com` and `http://todo.com/api` are received by a **Load Balancer** or **Proxy Server**. These requests are then directed to an **Ingress Controller Pod**. The Ingress Controller manages two **Service** objects: `todo-ui-service` and `todo-api-service`, both located on **Node 1**. The `todo-ui-service` is associated with port `3001` and the `todo-api-service` with port `8080`. Each service is represented by a blue box containing two **Pod** icons. Dashed arrows indicate the connection between the services and their respective pods.

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: todo-ingress-path-based
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /$1
spec:
  rules:
  - host: todo.com
    http:
      paths:
      - path: /?(.*)
        pathType: Prefix
        backend:
          service:
            name: todo-ui-service
            port:
              number: 3001
      - path: /api/?(.*)
        pathType: Prefix
        backend:
          service:
            name: todo-api-service
            port:
              number: 8080
```

Node 1

Cluster

PLAY

SUBSCRIBE

Ingress Service

› oc get ing

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
nginx-ingress	<none>	doc-nginx.apps.ocp-setu.waltoninstitute.ie	router-default.apps.ocp-setu.waltoninstitute.ie	80	33m

› oc describe ing/nginx-ingress

Name: nginx-ingress

Labels: <none>

Address: router-default.apps.ocp-setu.waltoninstitute.ie

Rules:

Host	Path	Backends
------	------	----------

doc-nginx.apps.ocp-setu.waltoninstitute.ie

/ nginx-service:8081 (10.129.3.3:80,10.130.1.247:80)

Volumes

- Problem 1: When a pod crashes, its data is lost.
 - The same applies to a node crash.
- Problem 2: The pods in a set cannot share data.
- Use volumes to overcome the above problems.
- A volume is a directory containing data that is accessible to a pod.
- Three types of volumes:
 1. emptydir – data outside the container, but private to a pod. Development only.
 2. hostPath – data outside the pod, but private to a node. Development only.
 3. Persistent Volumes – data outside the cluster and accessible by a pod set (cross-node).

Persistent Volumes

- Three components (Kubernetes resources):
 1. Persistent Volumes (PV)
 2. Persistent Volumes claims (PVC)
 3. Storage Classes
 - All are represented as cluster resources.
- A *Persistent Volume (PV)* is a piece of storage that has been provisioned by an administrator or dynamically provisioned.
 - Created with a YAML manifest.
- Storage Classes – types of storage devices, e.g. AWS EBS, NFS server, etc.
 - A PV specifies the storage class, size, access mode, etc.

Persistent Volumes

- A persistent volume claim (PVC) is a declaration by a pod that it will require storage.
 - A PVC is bound to a PV by matching the storage class, size, etc
 - The PV is mounted to a container path.
- Pod(s) → PVC → PV
- PV access modes:
 - ReadWriteMany – volume can be mounted by many nodes.
 - ReadWriteOnce – volume can be mounted by one node only. All the pods must be running on one worker node.
 - ReadMany / ReadOnce
 - ReadWriteOncePod

Application Configuration

- To use an application in different environments (dev, staging, prod), we should externalize its configuration, i.e. Reusability.
- Techniques:
 1. Command line arguments
 2. Configuration files
 3. Environment variables
- Similar techniques with Docker images:
 - Build-time arguments accessed in Dockerfile.
 - Runtime – environment variables used by container.
 - Avoids rebuild an image for each environment

Kubernetes – Pod configuration

- Avoid hard-coding pod configuration data in the deployment manifest.
- K8s has two special “volumes” for externalising pod configuring pod data:
 1. ConfigMaps.
 2. Secrets.

Kubernetes – ConfigMap

- ConfigMap is a K8s resource that lets you store config data that can be reused across applications.
- Config data can be:
 1. Simple key/value pairs, or
 2. An entire configuration file – use YAML | tp denote a literal block, i.e. the value is a multiline string and newline characters are preserved.
- Limited to 1 MB in size.
- Set immutable property to true to prevent updates.
 - Must delete and recreate to affect a change in configuration.
 - For mutable ConfigMaps, affected pods may need to be restarted after data changes.

Kubernetes – Secrets

- For confidential configuration data, use Secrets instead of ConfigMaps.
- Manifest files are identical to ConfigMaps.
- Data values must be encoded in BASE64.

```
$ echo -n plaintext | base64
```