



# Kubernetes (K8s)

# Background.

- An open-source system for automating the deployment, scaling and management of containerized applications.
  - A container orchestration platform.
- Originally announced by Google in 2014.
- In 2015, Google donated it to the Cloud Native Computing Foundation (CNCF).
- Kubernetes is the Greek word for helmsman – a person that steers a ship that carries containers of goods.

# The problems it solves.

- Problem 1: An application is comprised of multiple containers running on a cluster of nodes (on-premise server, Cloud VMs e.g. AWS EC2). A container or an entire node crashes.
  - K8s provides monitoring and self-healing, to ensure high availability.
- Problem 2: A container experiences a spike in traffic resulting in increased latency.
  - K8s provides auto-scaling and load balancing. It replicates the container on the node and/or across the cluster of nodes and load balances the requests between them. It scales up and down based on demand.

# The problems it solves.

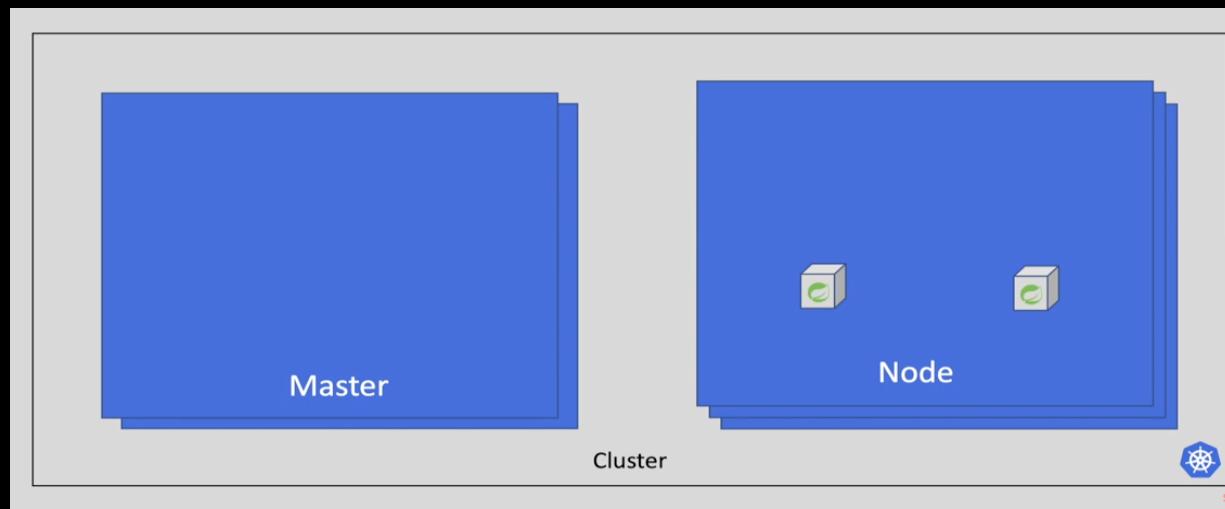
- Problem 3: You regularly upgrade an image and want to apply the change to a fleet of running containers without effecting downtime.
  - K8s provides rolling deployments - delete and restart each container, one by one rather than all at once.
  - Also supports canary deployments.
- Problem 4: I want to upgrade my entire application from v1 to v2.
  - K8s provides automatic rollout and rollback.

# The problems it solves.

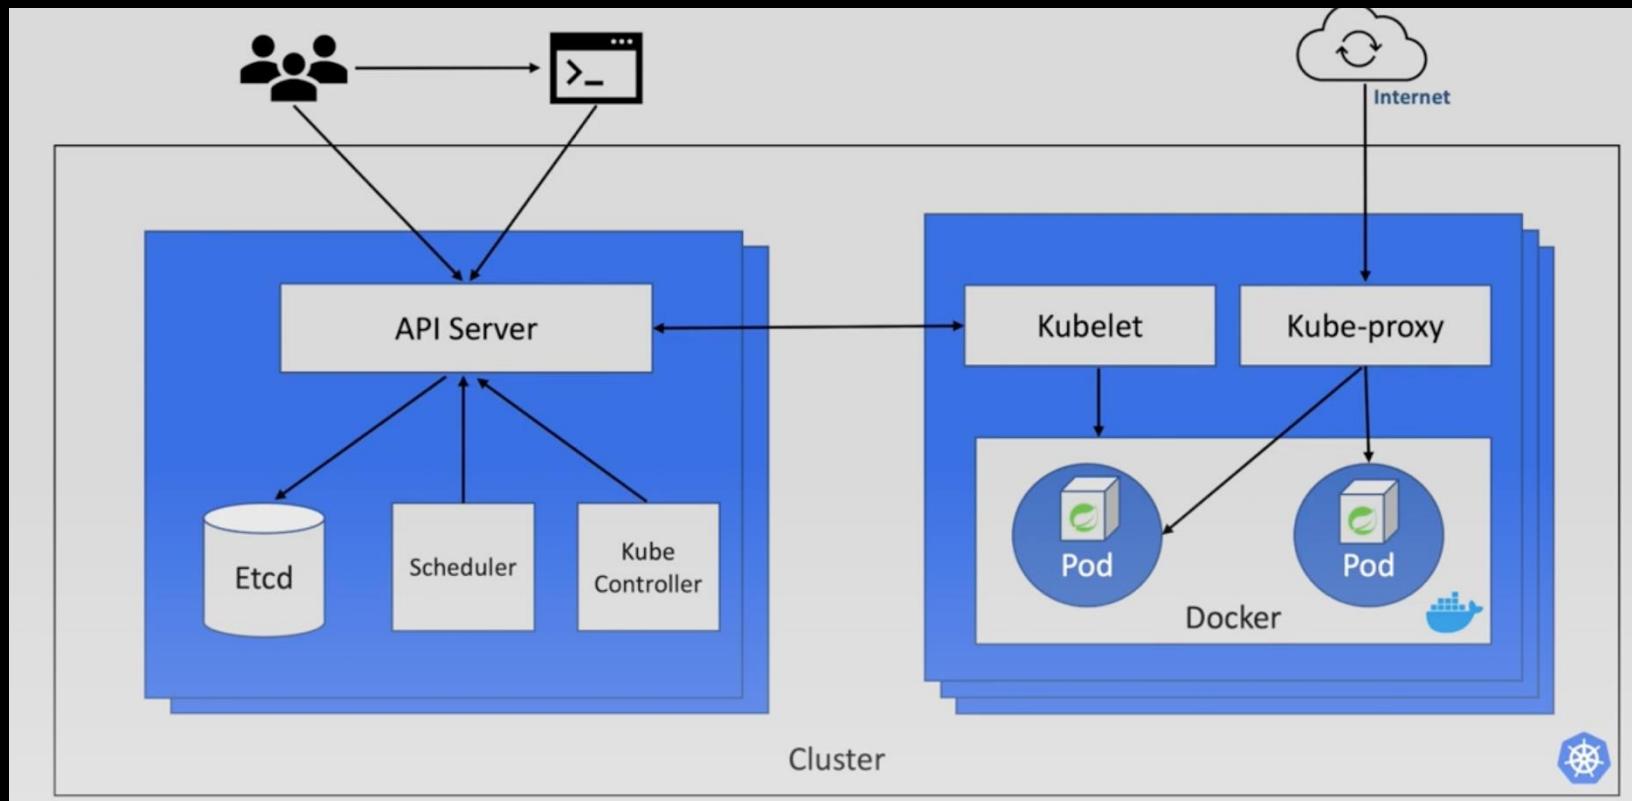
- **Problem 5: I want to reconfigure my application's without rebuilding images.**
- K8s provides secrets and configuration mappings to safeguard sensitive data and avoid unnecessary rebuilding of images.

# K8s Architecture

- K8s is installed over a set of nodes (VMs)
- Worker nodes host containers - the Data plane
- Master nodes (Control plane) manages the workers – More than one master for fault tolerance and high availability.
- The set of masters and workers is called a cluster

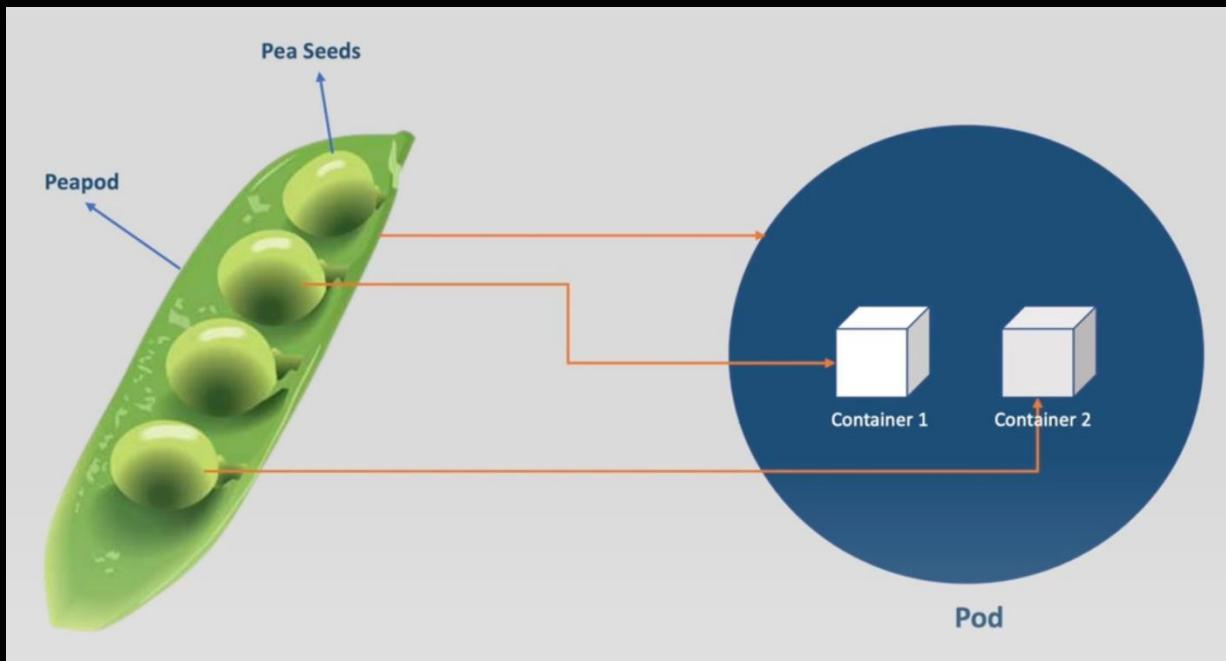


# K8s Architecture




# All about Pods

- The smallest unit of deployment is a pod.
- It encapsulates a group of containers, that share the same localhost network and storage.
- A pod is assigned a private IP address, but it's dynamic.



# All about Pods

- When a pair of containers are tightly coupled, they must reside on the same worker node. The pod construct ensures this is satisfied.



# Pod basics

- The containers in a pod share the same network (localhost), but use different ports.
- Deleting a pod will delete all its containers.
- To scale an app, we replicate its pod, where each one has its own group of containers.
- IP addressing:
  - A pod is assigned a private IP address.
  - On scale up, each pod replica is assigned its own IP.
  - Pod IPs are dynamic.

# Pod basics

```
> oc run nginx-pod --image=nginx:latest  
pod/nginx-pod created
```

```
> oc get pods  
NAME      READY  STATUS    RESTARTS   AGE  
nginx-pod  1/1    Running  0          12s
```

- Creating Kubernetes resources (e.g. pods) from the command line is cumbersome; instead, we use declarative code files for better maintainability.
  - Files are termed manifests – yaml or json options.

# Pod basics - Manifest files.

```
› oc apply -f 01-nginx-pod.yaml  
pod/nginx-pod1 created
```

```
› oc get pods  
NAME      READY  STATUS   RESTARTS  AGE  
nginx-pod  1/1    Running  0        3d5h  
nginx-pod1  1/1    Running  0        16s
```

```
› oc get pods -l team=integration  (Pod labels)  
NAME      READY  STATUS   RESTARTS  AGE  
nginx-pod1  1/1    Running  0        2m7s
```

# Pod basics - Attributes

- The etcd stores lots of attributes about pods

```
> oc get pod nginx-pod1 -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE
READINESS GATES							
nginx-pod1	1/1	Running	0	13m	10.130.1.156	worker1	<none>

- Use describe command to get all etcd info:

```
> oc describe pod nginx-pod1
```

- Response also shows the ‘events’ that occurred for a pod.

# Pod basics - Debugging

- Three main options:

1. Port-forwarding is possible with cluster access:

› oc port-forward nginx-pod1 3000:80

Forwarding from 127.0.0.1:3000 -> 80

Forwarding from [::1]:3000 -> 80

2. Check container's logs:

› oc logs nginx-pod1

3. Open a terminal shell inside the container:

› oc exec -it nginx-pod1 -- /bin/bash

root@nginx-pod1:/#

# ReplicaSets

- For high availability of an app (pod), K8s can create multiple copies in the cluster, called replica sets.

```
> oc apply -f 02-nginx-replica.yaml  
replicaset.apps/nginx-replica created
```

```
> oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replica-dmfvn	0/1	ContainerCreating	0	7s
nginx-replica-wc7j9	0/1	ContainerCreating	0	7s

```
> oc get rs (rs - replicaset)
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-replica	2	2	2	14s

# Self-healing

```
> oc delete pod nginx-replica-dmfvn
```

```
pod "nginx-replica-dmfvn" deleted
```

```
> oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replica-kkhs2	0/1	ContainerCreating	0	3s
nginx-replica-wc7j9	1/1	Running	0	7m36s

- The new pod's IP address is different to the one it replaced.
- A replica set's `spec.selector` property determines the pods it controls.
  - A pod's labels match the replica set's selector

# Deployments

- Deployments provide extra features over replica sets:
  - Rollout and rollback;
  - Deployment strategy.
  - Etc
- It automatically creates a replica set resource.
- The resource hierarchy:
  - A deployment manages a replica set.
  - A replica set manages a set of pods, based on matching the replica set's selectors with a pod's labels.
  - A pod manages its containers.

# Deployments

```
› oc apply -f 03-nginx-deployment.yaml
```

```
deployment.apps/nginx-depl created
```

```
› oc get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-depl	1/2	2	1	16s

```
› oc get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-depl-7569b665c5	2	2	2	20s

```
› oc get po
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-depl-7569b665c5-dvxrj	1/1	Running	0	25s
nginx-depl-7569b665c5-nd77g	1/1	Running	0	25s

# Deployment Updates

- Each Deployment update creates a new replica set

```
> oc apply -f 03-nginx-deployment.yaml. (After update)
```

```
deployment.apps/nginx-depl configured
```

```
> oc get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-depl-588f77bb97	2	2	1	15s
nginx-depl-7569b665c5	1	1	1	16m

```
> oc get rs
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-depl-588f77bb97	2	2	2	30s
nginx-depl-7569b665c5	0	0	0	16m

# Deployment Rollout

```
> oc apply -f 03-nginx-deployment.yaml (1st deploy)
```

```
deployment.apps/nginx-depl created
```

```
> oc rollout history deployment nginx-depl
```

```
deployment.apps/nginx-depl
```

```
REVISION CHANGE-CAUSE
```

```
1 Initial deployment
```

- Suppose we edit and apply the YAML twice:

```
> oc rollout history deployment nginx-depl
```

```
deployment.apps/nginx-depl
```

```
REVISION CHANGE-CAUSE
```

```
1 Initial deployment
```

```
2 Change to Nginx 1.21.5
```

```
3 Change to port 82
```

# Deployment Rollback


- Rollback – Undo a Deployment rollout.

```
> oc rollout undo deployment nginx-depl  
    --to-revision=2  
deployment.apps/nginx-depl rolled back
```

```
> oc rollout history deployment nginx-depl  
deployment.apps/nginx-depl  
REVISION  CHANGE-CAUSE  
1        Initial deployment  
3        Change to port 82  
4        Change to Nginx 1.21.5
```

# Services

- How do we communicate with a pod? i.e. other pods, external services, or end-user.
- Pods are assigned IP addresses, BUT these are NOT static and are not accessible outside the cluster.
  - Most rollouts and rollback cause pods to be recreated, resulting in new IP address assignments.
- Solution: Services



# Services

- Services abstract a set of pods (typically a replica set).
  - It's assigned an IP on creation, which is static.
  - Unlike pods, a service is not assigned to a node.
  - Services load balance the workload between its pods.
  - Service Types:
    1. ClusterIP (Accessible inside the cluster only)
    2. NodePort
    3. LoadBalancer
    4. Ingress

# Service creation

```
› oc apply -f 05-nginx-service.yaml
```

```
service/nginx-service created
```

```
› oc get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx-service	ClusterIP	172.30.250.186	<none>		
		8081/TCP	40s		

```
› oc get endpoints nginx-service
```

NAME	ENDPOINTS	AGE
nginx-service	10.130.0.75:80,10.131.1.91:80	2m10s

- These are the IPs (and ports) of its associated pods

# Service access

- Service IPs are private (only accessible inside the cluster).

```
› curl 172.30.250.186:8081
```

```
^C
```

```
› oc exec -it nginx-depl-7569b665c5-2mfzg -- sh
```

```
# curl 172.30.250.186:8081
```

```
<!DOCTYPE html><html>..... </html>
```

```
# curl nginx-service:8081
```

```
<!DOCTYPE html><html>.....</html>
```

```
#
```

- Pods can access services by their service name. due to K8s support for service discovery.
- Port-forwarding also allowed

# Services – Load balancing

- A service load balances the traffic to the pods

```
> oc exec -it nginx-depl-7569b665c5-2mfzg -- sh
```

```
# i=1
```

```
while [ $i -le 20 ]
```

```
do
```

```
curl nginx-service:8081;
```

```
i=$(( $i + 1 ))
```

```
done
```

- Monitor the pods by streaming their logs:

```
> oc logs nginx-depl-7569b665c5-2mfzg -f
```

# Other Service Types

- Unlike ClusterIP, NodePort and LoadBalancer types allow external access (outside the cluster), but both have deficiencies.
-

# Volumes

- Problem 1: When a pod crashes, its data is lost.
  - The same applies to a node crash.
- Problem 2: The pods in a set cannot share data.
- Use volumes to overcome the above problems.
- A volume is a directory containing data that is accessible.
- Three types of volumes:
  1. emptydir – data outside the container, but private to a pod
  2. hostPath – data outside the pod. but private to a node
  3. Persistent Volumes – data outside the cluster and accessible by a pod set

# Persistent Volumes

- Three components (Kubernetes resources):
  1. Persistent Volumes (PV)
  2. Persistent Volumes claims (PVC)
  3. Storage Classes
    - All are represented as cluster resources.
- A *Persistent Volume* (PV) is a piece of storage that has been provisioned by an administrator or dynamically provisioned.
  - Created with a YAML manifest.
- Storage Classes – types of storage devices, e.g. AWS EBS, NFS server, etc
  - A PV specifies the storage class, size, access mode, etc.

# Persistent Volumes

- A persistent volume claim (PVC) is a declaration by a pod that it will require storage.
  - A PVC is bound to a PV by matching class, size, etc
  - The PV is mounted to a container path.
- Pod(s) → PVC → PV
- PV access modes
  - ReadWriteMany – volume can be mounted by many nodes
  - ReadWriteOnce – volume can be mounted by one node only. All the pods must be running on one worker node.
  - ReadMany / ReadOnce
  - ReadWriteOncePod