



Kubernetes (K8s)

# Background.

- **An open-source system for automating the deployment, scaling and management of containerized applications.**
  - **A container orchestration platform.**
- **Originally announced by Google in 2014.**
- **In 2015, Google donated it to the Cloud Native Computing Foundation (CNCF).**
- **Kubernetes is the Greek word for helmsman – a person that steers a ship that carries containers of goods.**

# The problems it solves.

- **Problem 1: An application is comprised of multiple containers running on a cluster of nodes (on-premise server, Cloud VMs like AWS EC2). A container or an entire node crashes.**
  - **K8s provides monitoring and self-healing, to ensure high availability.**
- **Problem 2: A container experiences a spike in traffic resulting in increased latency.**
  - **K8s provides auto-scaling and load balancing. It replicates the container on the node and/or across the cluster of nodes and load balances the requests between them. It scales up and down based on demand.**

# The problems it solves.


- **Problem 3: You regularly upgrade an image and want to apply the change to a fleet of running containers without effecting downtime.**
  - **K8s provides rolling deployments - delete and restart each container, one by one rather than all at once.**
  - **Also supports canary deployments.**
- **Problem 4: I want to upgrade my entire application from v1 to v2.**
  - **K8s provides automatic rollout and rollback.**

# The problems it solves.


- **Problem 5: I want to reconfigure my application's without rebuilding images.**
- **K8s provides secrets and configuration mappings to safeguard sensitive data and avoid unnecessary rebuilding of images.**

# K8s Architecture

- K8s is installed over a set of nodes (VMs)
- Worker nodes host containers - the Data plane
- Master nodes (Control plane) manages the workers – More than one master for fault tolerance and high availability.
- The set of masters and workers is called a cluster




# K8s Architecture



# All about Pods

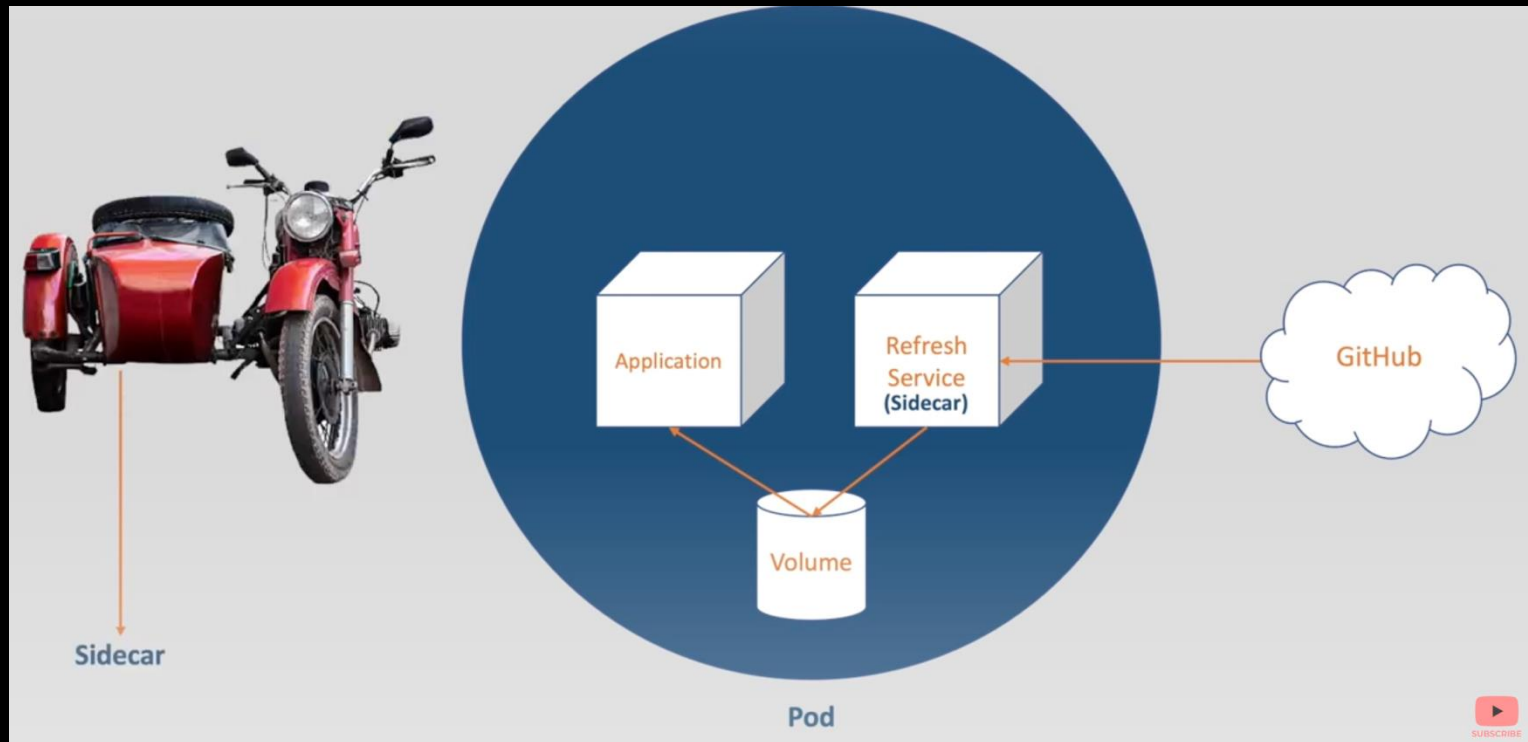
- The smallest unit of deployment is a pod.
- It encapsulates a group of containers, that share the same localhost network and storage.
- A pod is assigned a private IP address, but it's dynamic.





# All about Pods

- When a pair of containers are tightly coupled, they must reside on the same worker node. The pod construct ensures this is satisfied.



# Pod basics

- The containers in a pod share the same network (localhost), but use different ports.
- Deleting a pod will delete all its containers.
- To scale an app, we replicate its pod, where each one has its own group of containers.
- IP addressing:
  - A pod is assigned a private IP address.
  - On scale up, each pod replica is assigned its own IP.
  - Pod IPs are dynamic.

# Pod basics

```
➤ oc run nginx-pod --image=nginx:latest  
pod/nginx-pod created
```

```
➤ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-pod	1/1	Running	0	12s

- Creating Kubernetes resources (e.g. pods) from the command line is cumbersome; instead, we use declarative code files for better maintainability.
  - Files are termed manifests – yaml or json options.

# Pod basics - Manifest files.

➤ oc apply -f 01-nginx-pod.yaml

pod/nginx-pod1 created

➤ oc get pods

NAME	READY	STATUS	RESTARTS	AGE
nginx-pod	1/1	Running	0	3d5h
nginx-pod1	1/1	Running	0	16s

➤ oc get pods -l team=integration (Pod labels)

NAME	READY	STATUS	RESTARTS	AGE
nginx-pod1	1/1	Running	0	2m7s

# Pod basics - Attributes

- The etcd stores lots of attributes about pods

➤ oc get pod nginx-pod1 -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE
READINESS GATES								
nginx-pod1	1/1	Running	0	13m	10.130.1.156	worker1	<none>	<none>

- Use describe command to get all etcd info:
  - oc describe pod nginx-pod1
    - Response also shows the ‘events’ that occurred for a pod.

# Pod basics - Debugging

- Three main options:

1. Port-forwarding is possible with cluster access:

➤ `oc port-forward nginx-pod1 3000:80`

Forwarding from 127.0.0.1:3000 -> 80

Forwarding from [::1]:3000 -> 80

2. Check container's logs:

➤ `oc logs nginx-pod1`

3. Open a terminal shell inside the container:

➤ `oc exec -it nginx-pod1 -- /bin/bash`

`root@nginx-pod1:/#`

# ReplicaSets

- For high availability of an app (pod), K8s can create multiple copies in the cluster, called replica sets.

```
➤ oc apply -f 02-nginx-replica.yaml
```

```
replicaset.apps/nginx-replica created
```

```
➤ oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replica-dmfvn	0/1	ContainerCreating	0	7s
nginx-replica-wc7j9	0/1	ContainerCreating	0	7s

```
➤ oc get rs      (rs - replicaset)
```

NAME	DESIRED	CURRENT	READY	AGE
nginx-replica	2	2	2	14s

# Self-healing

```
> oc delete pod nginx-replica-dmfvn
```

```
pod "nginx-replica-dmfvn" deleted
```

```
> oc get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-replica-kkhs2	0/1	ContainerCreating	0	3s
nginx-replica-wc7j9	1/1	Running	0	7m36s

- The new pod's IP address is different to the one it replaced.
- A replica set's spec.selector property determines the pods it controls.
  - A pod's labels match the replica set's selector



To be continued ....